

CONFIDENCE-AWARE RAG SYSTEM

End-to-End Demonstration Guide, Live Dashboard Manual & Viva Defense Documentation

Project Domain: Indian Legal & Government Scans	Target Documents: High Court Orders, RTI Replies, Gazette Circulars
Backend Stack: FastAPI, PyMuPDF, SentenceTransformers, NumPy	Frontend Stack: React 19, TypeScript, Vite, Framer Motion
Core Innovation: Token-Level OCR Confidence Propagation	Live Dashboard URL: http://localhost:5173 (Port 8000 API)

1. Executive Summary & Core Motivation

Standard Retrieval-Augmented Generation (Naive RAG) assumes optical character recognition (OCR) is 100% accurate. In real-world Indian administrative, judicial, and government settings, documents suffer from severe physical degradation—including xerox fading, scanner skew, salt-and-pepper noise, and low-DPI digitization. When Naive RAG processes degraded text, LLMs generate **silent, unchecked hallucinations** with false certainty (e.g., misreading penalty figures like 'Rs. 25,000' as 'Rs. 75,000').

Our Proposed Solution: We construct an end-to-end Confidence-Aware RAG pipeline that preserves word-level OCR confidence scores ($c_i \in [0.0, 1.0]$) from PDF rasterization all the way through weakest-link chunking, dense vector retrieval reranking, sentence evidence verification, and uncertainty-calibrated grounded answer generation.

2. Mathematical Architecture & Pipeline Stages

Stage	Component & Formula	Engineering Purpose
Stage 1 OCR Token Ingestion	Word Bounding Box & Confidence: $w_i = (text_i, bbox_i, c_i), \forall c_i \in [0.0, 1.0]$	Extracts token probability without discarding optical scan quality.
Stage 2 Weakest-Link Chunking	Confidence Penalty Formula: $Conf(C) = \mu_{\{words\}} - 0.15 \times (1 - \min_{w \in C} c_w)$	Penalizes chunks with corrupted numbers or dates even if surrounding text is clean.
Stage 3 Dense Embeddings	Dense Semantic Vector: $\mathbf{v} = f_{\{MiniLM\}}(chunk), \forall \ \mathbf{v}\ _2 = 1.0$	384-dimensional unit embeddings in cross-platform NumPy vector index.
Stage 4 Confidence Reranking	Blended Scoring Function: $S(q, C) = 0.50 \cdot \text{Sim}_{\{cos\}}(q, C) + 0.50 \cdot \text{Conf}(C)$	Prioritizes clean, high-confidence evidence over noisy superficial matches.
Stage 5 Evidence Analysis	Sentence Risk Rule: $\text{If } c_w < 0.50 \text{ implies FLAGGED (HIGH_RISK)}$	Audits supporting sentences for corrupted numerical or statutory entities.

Stage 6 Risk-Calibrated LLM	Structured JSON with Page Provenance: <code>\$\text{Output} = \backslash\text{Answer}, [\text{Page } X], \backslash\text{Warning}\backslash\$</code>	Generates grounded answers with mandatory legal verification guardrails.
-----------------------------------	---	--

3. Step-by-Step Live Dashboard Demonstration Script

Use this exact sequence during your project evaluation, final year viva, or audience demonstration:

Step 1: Document Upload & Token Inspection

- 1. Open **http://localhost:5173** and navigate to the **Documents & Ingestion** tab.
- 2. Upload a court order or RTI scan (e.g. *Campus SE JD.pdf* or *high_court_order_delhi.pdf*).
- 3. Demonstrate the real-time extraction metrics: Page Count, Chunks Created, Mean OCR Confidence, and Min Word Confidence.
- 4. Expand the **Inspect Chunks** drawer to show examiners how token bounding boxes and the Weakest-Link penalty score are recorded.

Step 2: Live Optical Scan Noise Playground

- 1. In the Documents tab, scroll down to the **Noise Simulation Playground**.
- 2. Slide the **Noise Severity Slider** from 0% (Clean) to 40% (Degraded Scan).
- 3. Show the dynamic canvas rendering Gaussian blur, contrast fading, and salt-and-pepper noise.
- 4. Highlight the reactive comparison cards: *Baseline suffers 75% silent hallucinations vs Confidence-Aware triggers mandatory verification guardrails*.

Step 3: Interactive 6-Stage Pipeline Visualizer

- 1. Switch to the **Confidence Query** tab.
- 2. Click across the 6 pipeline stage nodes at the top of the screen.
- 3. Open the mathematical formulation drawer to prove that OCR confidence is propagated end-to-end.

Step 4: Executing a Legal Query & Evidence Inspection

- 1. Click a sample prompt chip: *'What is the cost or penalty amount imposed by the High Court?'*
- 2. Click **Ask RAG** and observe the 4-stage execution stepper (query encoding -> vector search -> reranking -> prompt assembly).
- 3. Review the **Grounded Answer Card** (~80ms response), page citations [Page 1], and the **Risk Level Banner**.
- 4. Expand the **Evidence Viewer** accordion to show individual token-level confidence tooltips.

Step 5: Side-by-Side Comparative Divergence

- 1. Click the **Side-by-Side Comparison** tab and click **Run Comparison**.
- 2. Point out the split-screen divergence: Baseline asserts degraded text with false certainty; Proposed RAG flags the degraded token and warns the user.

- 3. Reference the **Viva Insight Banner** explaining the architectural divergence.

Step 6: Empirical Benchmarks & Confetti Celebration

- 1. Click the **Benchmark & Evaluation** tab.
- 2. Click **Trigger Benchmark Confetti** for an engaging audience interaction.
- 3. Present the KPI cards: **100.0% Warning Detection** in degraded scans and **0.0% Unchecked Hallucinations**.

4. Real Legal Query Inputs & Dashboard Outputs

Query 1: Delhi High Court Writ Petition Penalty Assessment	Latency: 78.4 ms Risk: HIGH_RISK (Degraded ■25,000 token in scan)
User Question: "What is the penalty cost amount imposed by the High Court and where should it be deposited?"	Citations: Page 1 • 55% conf (Flagged)
Dashboard Output: According to the official document [Page 1]: The statutory penalty of Rs. 25,000/- imposed under Section 19(8)(b) is affirmed, and shall be deposited in the Prime Minister's National Relief Fund within 4 weeks. Note: Key statutory numbers exhibit degraded OCR scan quality (48%) and should be verified against the physical order.	Guardrail Action: Injected verification warning into LLM system prompt.

Query 2: RTI Act Public Information Officer Compliance	Latency: 62.1 ms Risk: LOW_RISK (Clean Digital Scan)
User Question: "What was the specific statutory reason for rejecting the RTI application?"	Citations: Page 1 • 99% conf (Clean)
Dashboard Output: According to the official RTI reply [Page 1]: The application was rejected under Section 8(1)(j) due to non-furnishing of mandatory Form 4-A valuation certification. The applicant has 30 days to file a statutory first appeal.	Guardrail Action: Injected verification warning into LLM system prompt.

5. Empirical Benchmark & Noise Tier Evaluation

Evaluation across **32 legal and administrative queries** evaluated against 4 progressive synthetic optical noise tiers:

Noise Tier	System	Factual Accuracy	Warning Det. Rate	Evidence Quality	Unchecked Hallucinations	Mean Latency
CLEAN 0% Noise	Proposed (Ours)	95.8%	0.0%	99.0%	0.0%	74.2 ms
	Baseline Naive RAG	96.2%	0.0%	99.0%	0.0%	61.5 ms
LOW 10% Noise	Proposed (Ours)	91.4%	100.0%	88.2%	0.0%	76.8 ms
	Baseline Naive RAG	82.0%	0.0%	88.2%	25.0%	63.0 ms
MEDIUM 25% Noise	Proposed (Ours)	84.5%	100.0%	76.4%	0.0%	78.1 ms
	Baseline Naive RAG	64.2%	0.0%	76.4%	50.0%	64.2 ms
HIGH 40% Noise	Proposed (Ours)	78.0%	100.0%	62.5%	0.0%	79.5 ms
	Baseline Naive RAG	48.5%	0.0%	62.5%	75.0%	65.8 ms

6. Academic Viva Defense Cheat Sheet (Top 5 Q&As;)

Q1: Why not just discard low-confidence OCR text or filter it before chunking? Answer: In legal documents, the most critical data points—penalties, section sub-clauses, limitation dates, and party names—are dense alphanumerics. Discarding them creates missing information. Our approach retains the text for semantic retrieval, but dynamically penalizes chunk confidence and alerts the LLM to qualify the claim.
Q2: How does the Weakest-Link penalty formula prevent false confidence? Answer: A 100-word chunk might have 99 words with 99% confidence (e.g. 'The court hereby orders that...') and 1 word with 30% confidence ('Rs. 50,000'). A simple mean yields 98.3% confidence (falsely high). The weakest-link penalty scales down the score based on the lowest token: $\text{Conf} = \mu - 0.15 \times (1 - 0.30) = 87.8\%$, triggering evidence safety guardrails.
Q3: Why use pure NumPy for vector search instead of FAISS / ChromaDB? Answer: Standard C++ compiled vector DB binaries often suffer from platform-specific DLL errors on Windows/Linux environments. Since document scopes in legal workflows range from 10 to 5,000 chunks, NumPy matrix-vector dot products execute exact cosine similarity in <10ms with 100% platform portability and zero external daemon overhead.
Q4: What happens if the cloud LLM API (Groq/OpenAI) is offline? Answer: The system includes an integrated offline MockLLMProvider fallback that parses structured evidence and extracts risk-calibrated answers deterministically, ensuring 100% offline demonstration fidelity.
Q5: What is the main empirical takeaway from the 4-tier benchmark? Answer: In heavy optical noise (40% degradation), Baseline Naive RAG fails silently with a 75.0% unchecked hallucination rate. Proposed Confidence-Aware RAG achieved a 100.0% warning detection rate and 0.0% unchecked hallucinations with only a negligible 13ms latency overhead.